

A

Seminar Report

Comparison and Performance Analysis of Standard and LVM Based Disk Parititioning

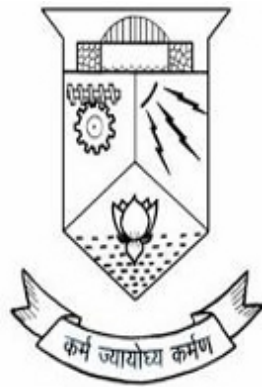
Submitted in partial fulfillment of the requirements for the Award of the Degree

of

Master of Computer Applications

of

APJ Abdul Kalam Technological University



Submitted by

ASWIN BABU K

RegNo: TVE16MCA14

Department of Computer Applications

COLLEGE OF ENGINEERING TRIVANDRUM

SEPTEMBER 2018

DEPARTMENT OF COMPUTER APPLICATIONS

COLLEGE OF ENGINEERING TRIVANDRUM



CERTIFICATE

*Certified that this Seminar report entitled, “**Comparison and Performance Analysis of Standard and LVM Based Disk Partitioning**” is the paper presented by “**Aswin Babu K**” (Reg No: **TVE16MCA14**) in partial fulfillment of the requirements for the award of the degree of Master of Computer Applications of APJ Abdul Kalam Technological University during the year 2018.*

Prof. Baby Sylal.

Prof. Jose T Joseph.

Co-ordinator

Head

of the Department

Acknowledgement

First and for most I thank **GOD** almighty and to my parents for the success of this seminar. I owe my profound gratitude and sincere thanks to everyone who shared their precious time and knowledge for the successful completion of my seminar.

I would like to thank **Dr. Jiji C.V**, Principal, College of Engineering Trivandrum, who helped me during the entire process of work.

I am extremely grateful to **Prof. Jose T Joseph**, HOD, Dept of Computer Applications, for providing me with best facilities and atmosphere for the creative work guidance and encouragement.

I would like to thank my coordinator, **Prof. Baby Sylva**, Dept of Computer Applications, who motivated me throughout the work of my seminar.

I profusely thank other Asst. Professors in the department and all other staffs of CET, for their guidance and inspirations throughout my course of study.

I owe my thanks to my friends and all others who have directly or indirectly helped me in the successful completion of this seminar. No words can express my humble gratitude to my beloved parents and relatives who have been guiding me in all walks of my journey.

Aswin Babu K

Abstract

In the course of twenty five years, Linux has gone from being a hobbyist toy Operating System run on the creator's 386 machine to an OS that runs everything from smart watches to Supercomputers. From managing Hard Disk drives having a few hundred megabytes of storage, current Linux systems are tasked with management of vast storage arrays, whose size can go up to tens of Petabytes. From the advent of modern storage devices, disk partitioning was seen as a necessity, as it allowed proper allocation of storage resources for various users and services of the system. But traditional disk partitioning has struggled to keep up with the exhilarating advancements in the field of storage. Logical Volume Manager or LVM is Linux's answer for easing the storage management. It does this by further abstracting the storage partitions allowing them to span multiple physical disks and to be resized even while being mounted. The seminar introduces standard and LVM based partitioning methods and then looks if LVM can be used as a replacement for the standard partitioning methods.

Contents

| | | |
|----------|-------------------------------------|-----------|
| 1 | Introduction to File Systems | 1 |
| 2 | Disk Partitioning | 2 |
| 3 | Partition table | 3 |
| 3.1 | Types | 3 |
| 3.2 | Disadvantages | 4 |
| 4 | Logical Volume Manager | 6 |
| 4.1 | Implementation | 6 |
| 4.2 | Snapshots | 7 |
| 4.3 | Striped Volume | 8 |
| 4.4 | Advantages | 8 |
| 4.5 | Disadvantages | 8 |
| 5 | RAID | 9 |
| 5.1 | RAID Striping | 9 |
| 5.2 | Disadvantages | 10 |
| 6 | LVM Performance | 11 |
| 7 | Conclusion | 13 |

List of Figures

| | | |
|-----|--|---|
| 3.1 | Typical Partition table on a physical disk | 4 |
| 4.1 | Typical LVM structure | 7 |

Chapter 1

Introduction to File Systems

Computers use secondary storage devices to permanently store data. To write the data, a computer uses one of the widely known specifications. This allows the computer to read back the data, once written. These specifications are more commonly known by the term File Systems. File Systems enforce a uniform method of storing and retrieving the data. Further, their popular nature allows the device to be read properly across Operating Systems. You must have come across at least a couple of file systems, if you have used a computer. Windows use NTFS as its file system. Linux people prefer to use ext4 or btrfs. Your USB drive most probably uses the FAT32 file system.

So when you copy a file to your hard disk, you are writing to the file system right? Actually, it is a bit more complicated than that. The computer does not simply create a file system on top of your physical device. To create a file system, you first need an abstraction, called a partition. Partition is a logical slice of your physical device. It acts as a container onto which a file system can be created. The partition maps to further abstracted versions of physical locations on the drive. Suffice to say, a partition does not necessarily have to occupy the whole storage device. There can be more than one partition and each partition can have different file systems, independent from the rest of the partitions.

The process of creating partitions to store data on a storage device is called disk partitioning. Apparently, ghosts of magnetic storage media are going to haunt us for a very long time. You can read more about disk partitioning in the next chapter.

Chapter 2

Disk Partitioning

The term disk partitioning might bring up memories of slicing a pie. But let me assure you that it is an entirely logical process. Disk partitioning makes a secondary storage device usable. It creates partitions, which are then formatted with a file system of our choice, and then used to store data. What if you just want a single partition that spans an entire drive? You can ofcourse do that. Still, the process you performed will be called partitioning.

It is very common to see systems with multiple partitions, and with a good reason. Multiple partitions allow your system to be more flexible. There are plethora of use cases for splitting up your disk, some of which are listed below.

- Separate personal information from Operating System files
- Separate boot information
- Run multiple operating systems
- Further partition personal data
- To have a faster partition for virtual memory

Never forget the fact that, each of these partitions are mere abstractions. They are logical divisions which start from a sector of the drive and ends on an other sector. If the partitions are purely logical, someone or something has to keep track of their locations right? this is the job of the partition table.

Chapter 3

Partition table

You have partitioned the device as per your heart's content. But as said before, the partitions are purely logical creatures. They are abstractions. So where does the computer store information on partitions? They can't be stored inside the Operating System, as the data is needed first to understand where the OS is stored. This is the job of the partition table. A partition table is usually stored in the beginning of the drive. Every piece of low level software, that tries to read data from the drive accesses the partition table first. Like the plethora of file systems thriving in the wild, there are quite a few partition table formats used by the leading platforms.

3.1 Types

Historically, the technology industry has taken great pride in keeping all the knowledge they have gained to themselves. This meant new companies, trying to bring some new product into the market often had to reinvent stuff that others already had.

As a result, there are quite a few types of partition tables used by leading platforms. Apple uses something called Apple Partition Map, BSD uses disklabel but the most commonly used of them all is the Master Boot Record (MBR) and the GUID (GPT) Partition Table.

If you are not the tinkering type and you own a decade old PC, your computer most probably uses MBR. MBR started its life in the early 80s and carries with it, all the limitations of the decade. It originally allowed only four partitions (which was later fixed) and does not support drives of capacities more than 2 TB. This led Intel to develop a modern replacement called

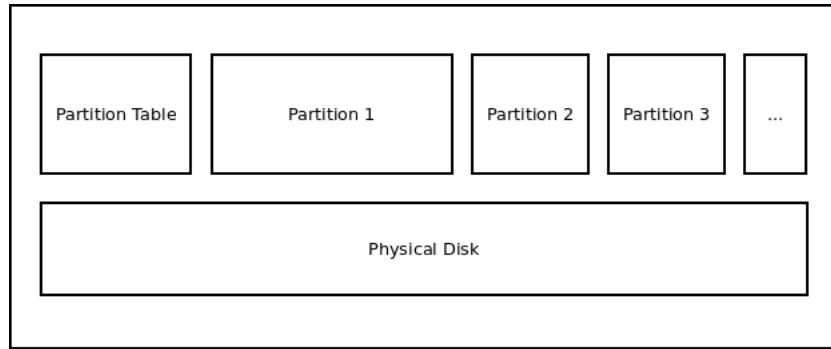


Figure 3.1: Typical Partition table on a physical disk

GPT. GPT is the partition table used by most of the modern PCs. It has overcome most of the limitations of the MBR.

So let's get this straight. You have the physical disk, on top of which, the partition table exists. The partition table stores information on all those imaginary partitions existing on the disk. The partitions are not physically etched on the disk. This should mean they are pretty flexible right? Actually, in reality, it is not that straightforward.

3.2 Disadvantages

Let us imagine that we have a 500 GB Hard Disk, in which 465 GB is the actual usable space (HDDs really do these nasty things). We split the drive into four partitions as follows:

- Partition 1 of size 50 GB for /
- Partition 2 of size 150 GB for /home
- Partition 3 of size 200 GB for Windows
- Partition 4 of size 65 GB for Backup

The drive is partitioned under the assumption that it will suffice for some time. But within weeks, the / partition starts filling up. Your Windows partition is virtually empty and you try to resize Windows so that you can get some space for / partition except that you can't. The / partition can only acquire the free space which is adjacent to it. And that by the way is the case

with each partition on the drive.

Further, if you want to modify the partition while the system is running, you are out of luck. This is disastrous if you think of a production server which is constantly serving the needs of clients. Most often, traditional partition tables force you to boot into a live environment, if you want to modify anything which the OS depends on.

Lets extend the production server concept bit further. Servers need lots of space. If your firm is doing something worthwhile, space will eventually run out and you will be forced to add more physical drives. Now you want the partition having all your critical data to span across both the drives. There actually exists a method called "Spanned Volume" on Windows which can do this but not before destroying your current partition first. So yes, traditional partitioning leaves you out in the cold when such situations arise. These are some of the problems that gets fixed, when you switch to LVM.

Chapter 4

Logical Volume Manager

If partitions are abstractions, Logical Volume Manager or LVM is an abstraction on top of abstraction. So far we have talked about partitions, which are logical structures on top of the physical disk. LVM extends this one step further by creating logical volumes, which are usable partitions, on top of existing partitions.

LVM is Linux's answer to all the problems discussed in the previous section, and it can do a lot more. In LVM's lingo, the partitions on your drive is called Physical Volume (PV). One or more physical volumes group together to form a Volume Group (VG). A volume group appears as a single continuous space no matter how many physical volumes lie underneath. LVM works by creating partitions called Logical Volumes (LV) on top of the volume group.

By virtue of definition, the logical volumes can span across physical drives, also they bring a number of goodies to the table, the most prominent being the ability to resize partitions on the fly. This means, switching into a live environment or even unmounting the partition is no longer necessary.

4.1 Implementation

The high level, logical nature of LVM means Linux cannot access it like the traditional partitions. To solve the issue, Linux uses a piece of code called Device Mapper. Device Mapper passes data from a virtual device such as a logical volume to the physical storage device. The snapshot functionality, about which you will read in the next section is implemented by the device mapper.

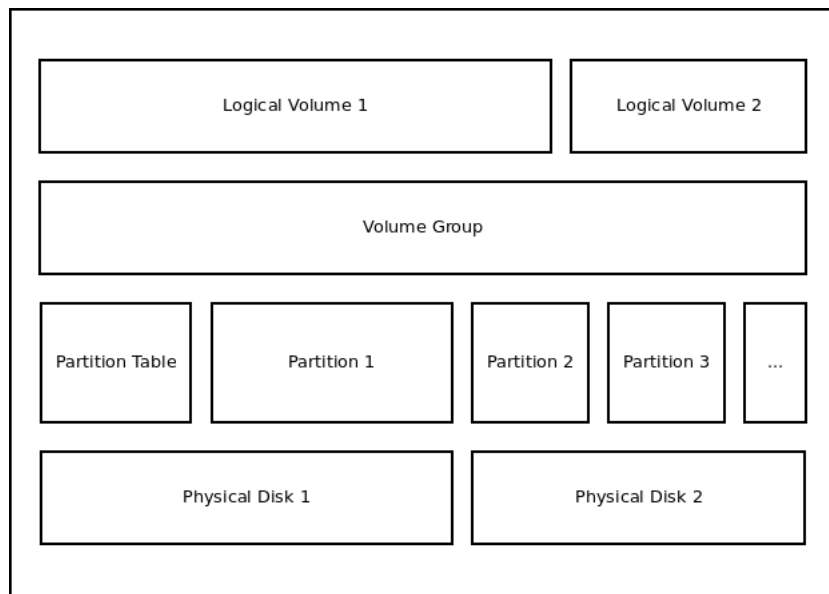


Figure 4.1: Typical LVM structure

The use of device mapper extends beyond implementing the LVM. It can be used for encrypting disks and simulating hardware behaviour.

4.2 Snapshots

Disks fail. Operating Systems crashes. Viruses can attack. All of these can take your precious data with them. Backing up data is an option, but what if you want to save the current state of your storage? This is where disk imaging comes into play. Disk imaging allows the creation of image of your entire disk. Normally, this is done by rebooting your computer and booting into a live environment such as CloneZilla. Which means, you guessed it, more downtime.

LVM provides a very handy feature called snapshots, which are like disk images, except that you can create one while your system is up. Snapshots allows you to create images of your entire storage which can be used to restore disks to the state at the time of creation of the snapshot. The advantage is obvious. You can continue to server your clients or do your work, while LVM creates a copy of your entire drive for future use. Once created, these images act like logical volumes, which can be mounted if required.

4.3 Striped Volume

Lets imagine you have two independant physical drives, which are capable of operating, well, independantly. That means, while one drive is doing its job, the one is free to do something else. Enter striped volume.

Striped volumes are logical volumes, that can span more than one physical drive. The advantage is obvious. You will reading and writing at twice the speed of a single drive. There are alternatives to this, such as using RAID. But LVM shines through due to ease of setting up the system.

4.4 Advantages

- Ability to resize partition without shutting down or even unmounting
- A single partition can span one or more physical drives
- The underlying disk drives can be of any size
- Take snapshot of entire storage while the system is up
- Striped volumes that can provide enhanced performance

4.5 Disadvantages

- LVM is Linux only and is not supported under Windows
- If a physical drive fails while using striped volume, data can be lost
- LVM cannot offer raw disk performance

Chapter 5

RAID

RAID stands for Redundant Array of Inexpensive (or Independant) Disks depending on whom you ask. RAID is a method for combining disk drives for increasing performance, reliability or both. RAID can operate differently according to mode in which it operates. There are seven standard modes numbered from 0 to 7. We are not going to discuss all of them as it exceeds the scope of this report. RAID provides a number of features that are similar to the ones provided by LVM, but differences exist, which are discussed in this chapter.

5.1 RAID Striping

Except RAID 1, all RAID levels support the creation of striped partitions. If RAID is implemented through a dedicated hardware card, the striping is transparent to the Operating System and as a side effect, performs considerably faster than an LVM striped volume.

When you create a striped volume, the main concern is the reliability of the drives working underneath. If a drive fails, it can make the data inaccessible or useless. RAID level 5 and up tackles this issue by using distributed parity. In simpler words, data correction mechanisms make sure that the data stays accesible even if one of the drives in the array fails. But the technology is not without its faults.

5.2 Disadvantages

First of all, RAID striping requires that all of the drives be of same capacity. If the computer initially had a 500 GB drive and you bought an additional 1 TB drive to do striping, congrats. Now you have two drives with the usable size of 500 GB. In other words, the extra 500 GB of your new HDD cannot be used. The usable size of each drive in a RAID array implementing striping is the size of the drive having the least size. So if your computer initially had a 500 GB drive, you have to keep buying drives of the capacity 500 GB or throw out the older drive. I call that expensive.

RAID is fast. But to be fast, your computer should have a dedicated RAID card, and these cards do not come cheap. Further, RAID 5 requires that the system have at least three drives to begin with. All this makes RAIDing the system uncomfortably out of reach of a normal consumer.

Even if you are ready to invest in a pricey RAID card and at least three HDDs, RAID still won't offer the flexibility of resizing the partitions provided by LVM. If you wish to extend the partition once a new drive is added, expect some downtime. Also, RAID does not offer snapshots. So if you value your data, you willing eventually end up implementing LVM on top of your pricey RAID setup.

Chapter 6

LVM Performance

When using LVM, all the abstraction layers mean the computer has to process a lot of code before hitting the bare metal. This inturn should lead to lower performance. But exactly how much performance is lost in this transition? Four independant researchers named Milivoj Bozic, Milos Djukic, Dragan Narancic and Istvan Pap decided to find out.

They installed 32-bit Linux of Kernel revision 3.14 on two identical single board computers. The only difference between these two setups was that, the first one had traditional partitioning, while the second one employed LVM. Both the traditional and LVM setup had two partitions, one of 10 GB and the other of the size 20 GB. Further, the LVM setup was tested in two configurations. First with version 1 of LVM and second with version 2 of LVM.

While the non-LVM setup took 19 seconds to boot the machine, the LVM and LVM2 configurations took 23 and 24 seconds respectively. On an average, the LVM setup takes a 20% performance hit, compared to the non-LVM setup. This is because the GRUB bootloader has to load an additional module allowing it to recognize the LVM volumes.

When it comes to file read/write performance, the performance hit is less pronounced. While the non-LVM configuration delivers 72 MB/s during the write operation of a 100 MB file, the same task delivers a performance of 67 MB/s on the LVM setup. The experiments proved that on average, the non-LVM setup was only 2% to 5% faster than the LVM setup.

Further, the team also tried to benchmark the configurations using a RAID setup. It was found that RAID setup takes a 15% performance hit. This is likely due to the onboard RAID controllers used for the research.

Chapter 7

Conclusion

Though the Logical Volume Manager is a fairly complex storage system, it has been engineered to provide performance figures similar to the ones provided by traditional partitioning.

LVM is easy to setup, is practically free and provides a number of features such as convenient backups through snapshots, easy resizing of partitions and improved performance using striped volumes while providing satisfactory performance. The study concluded that LVM is capable of being a viable replacement to traditional partitioning in both server and development environments.

Bibliography

- [1] Milivoj Boznic, Milos Djukic, Dragan Narancic, Istvan Pap, “*Comparison and Analysis of Standard disk Partitioning and LVM based disk Partitioning on Linux Systems*”, TELSIS 2017